



NGINX: Basics and Best Practices

NGINX

Agenda

- Introducing NGINX
- ADC Augment and Modernization
- Installing NGINX and NGINX Plus
- Essential files, commands, and directories
- Basic configurations
- Advanced configurations
- Monitoring and Logging
- Summary

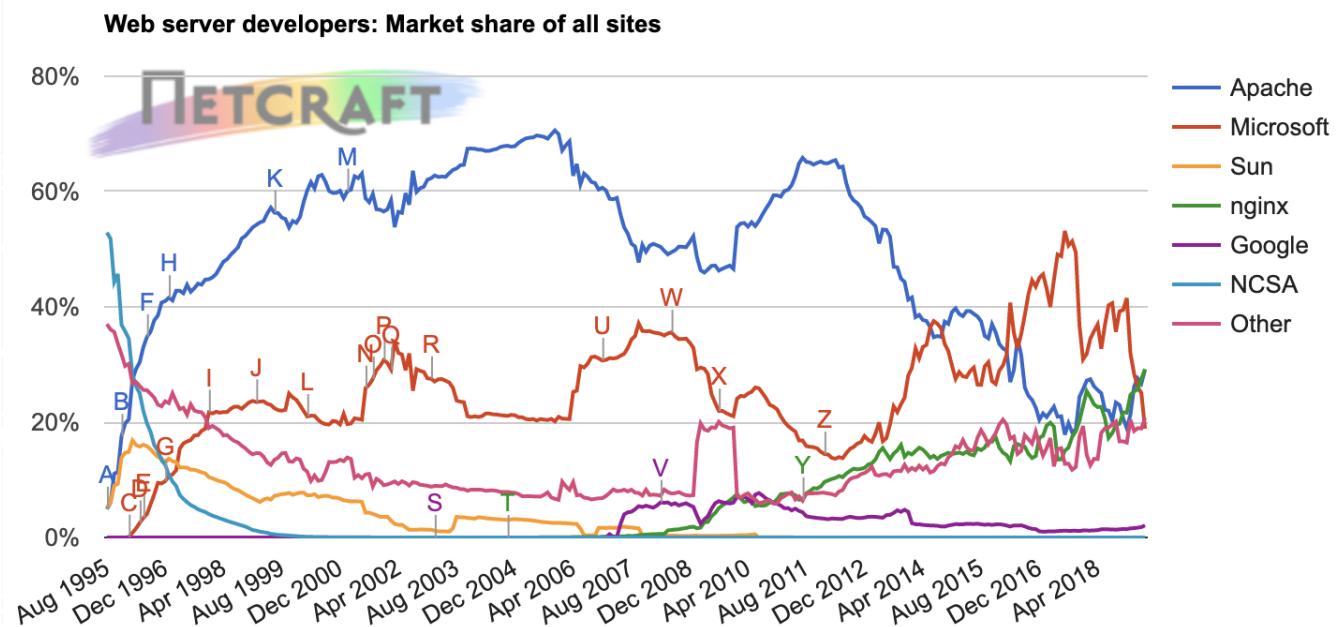




Introducing Nginx

**NGINX is the
most used
web server
on the
internet**

Source: [w3techs, May 2019](#)



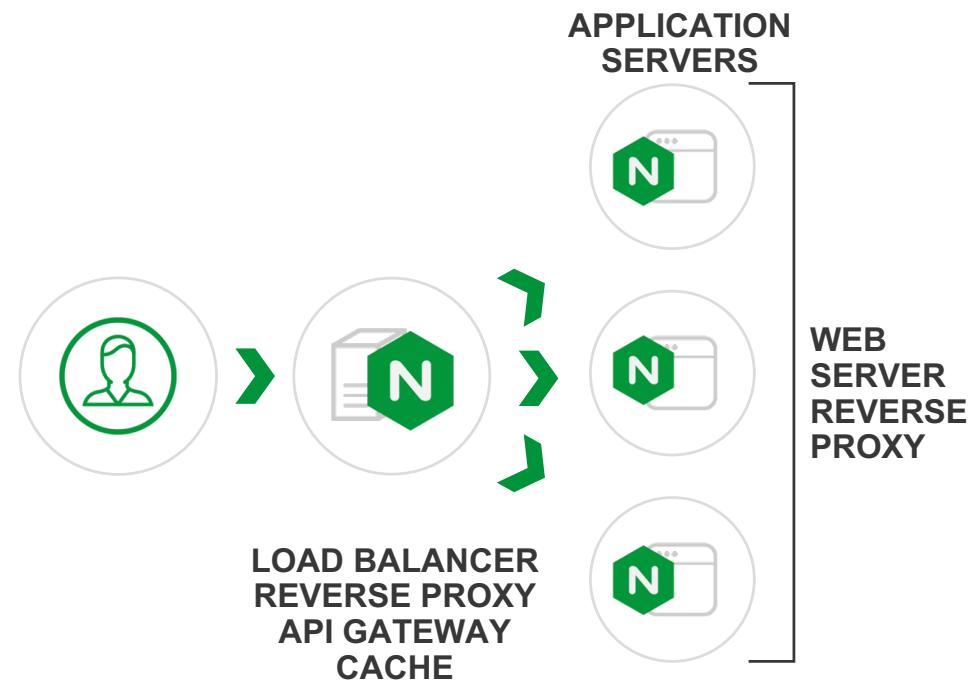
What is NGINX?

NGINX

- Basic load balancer
- Reverse Proxy and Web Server
- Content Cache
- SSL termination
- Rate limiting
- Basic authentication

NGINX PLUS

- Active health checks
- Session persistence
- DNS service discovery integration
- Cache-purging API
- JWT authentication and OpenID Connect
- Live Activity monitoring (100+ real time metrics)
- Dynamic Modules
- API for Dynamic reconfiguration, Cache-purge, key-value store
- High Availability, Cluster State syncmuch more.

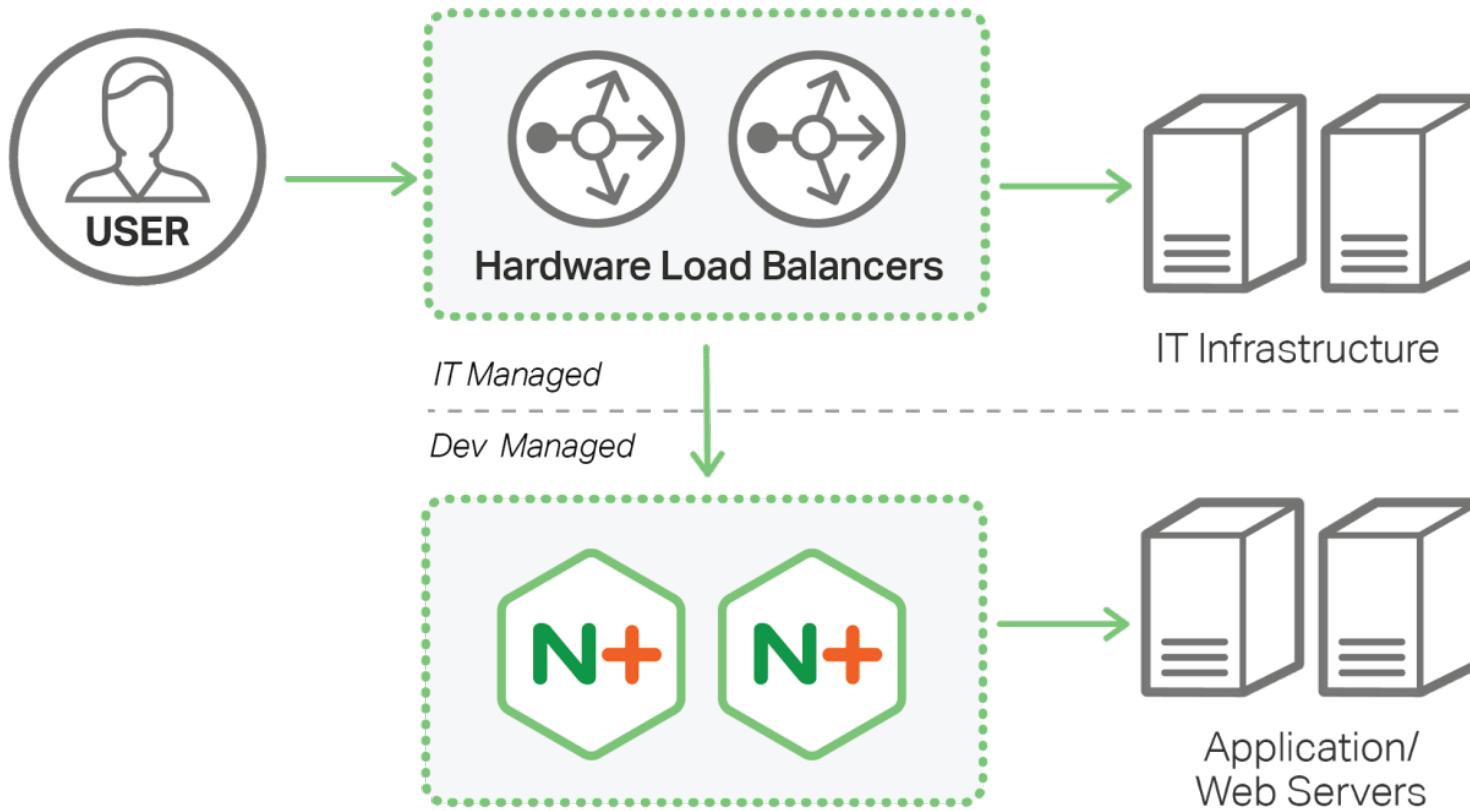




ADC Augment and Modernization

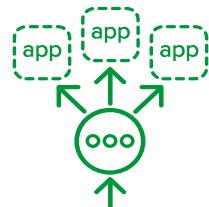
What's happening now

Traditional Application Infrastructure are being augmented



ADC Augment - key use cases

Key use cases



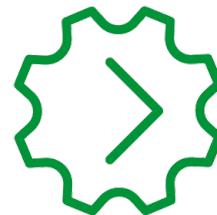
ADC Augment

Enhancing existing app environments



ADC for Multi-Cloud

Scale and Secure Apps across multi-cloud



API Management

End-to-end API lifecycle services

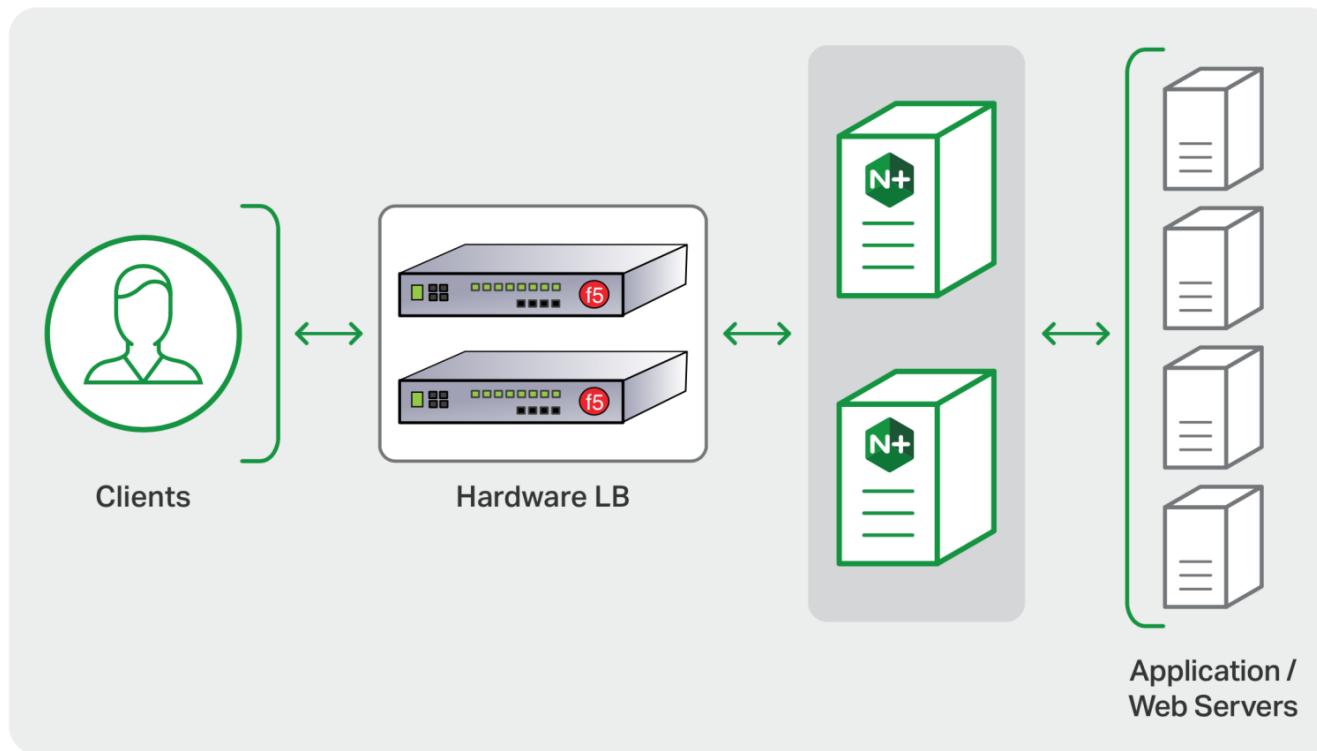


Kubernetes Integration

Flexible and scalable app services

1. Augment Traditional Load Balancers

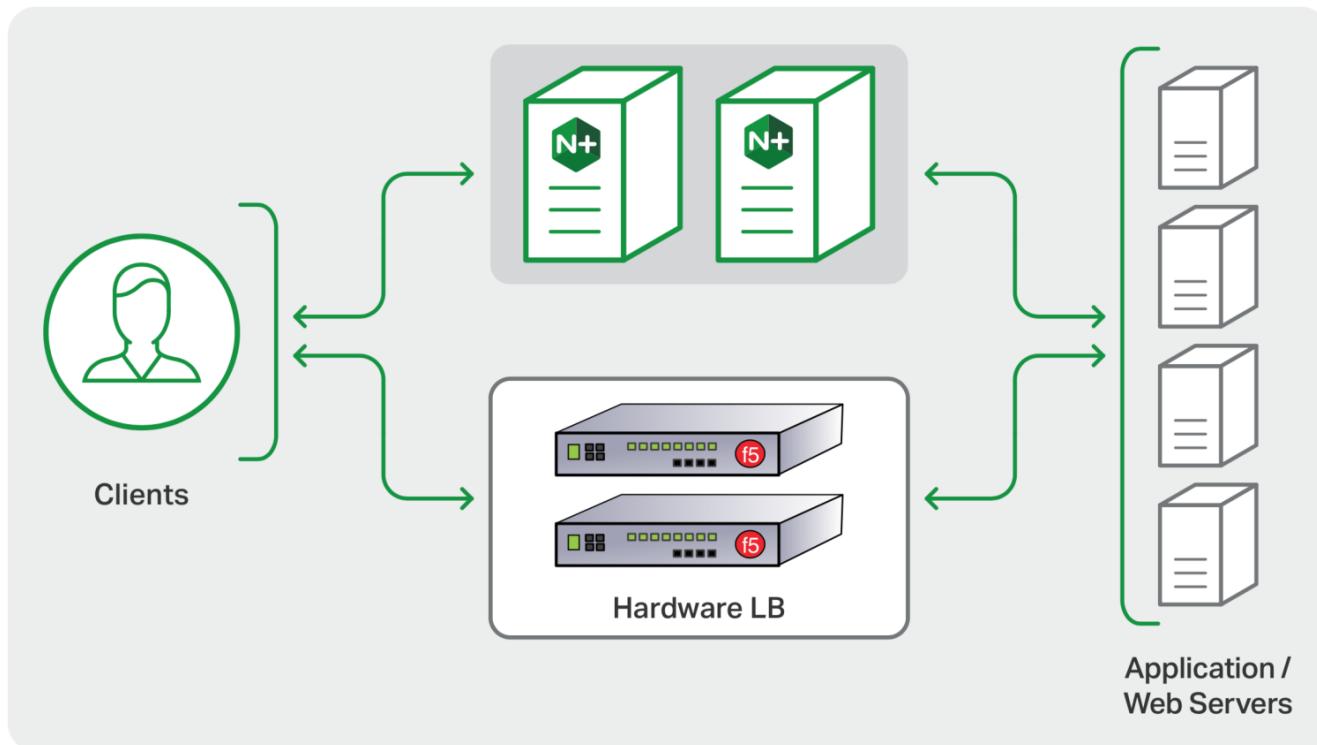
Traditional Application Infrastructure are being augmented



- Easiest way to introduce NGINX into your network
- Hardware layer 4 load balancer to NGINX
- Can start small with one application being behind NGINX and then expand

2. NGINX Alongside Hardware ADCs

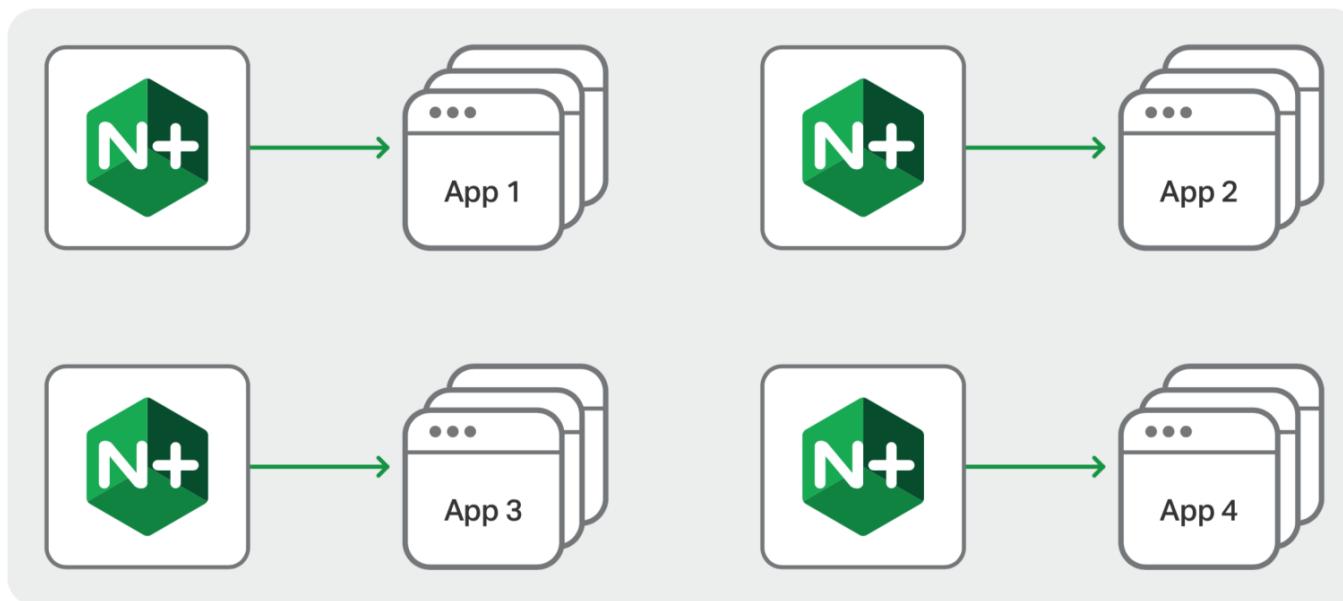
Offload or Migrate new application workloads



- Parallel NGINX deployment
- Good architecture if adopting public cloud while still keeping private datacenter
- Can also start small with one application being behind NGINX and then expand

3. Micro Load Balancers/Gateways

Legacy Hardware ADC replace to a application centric architecture



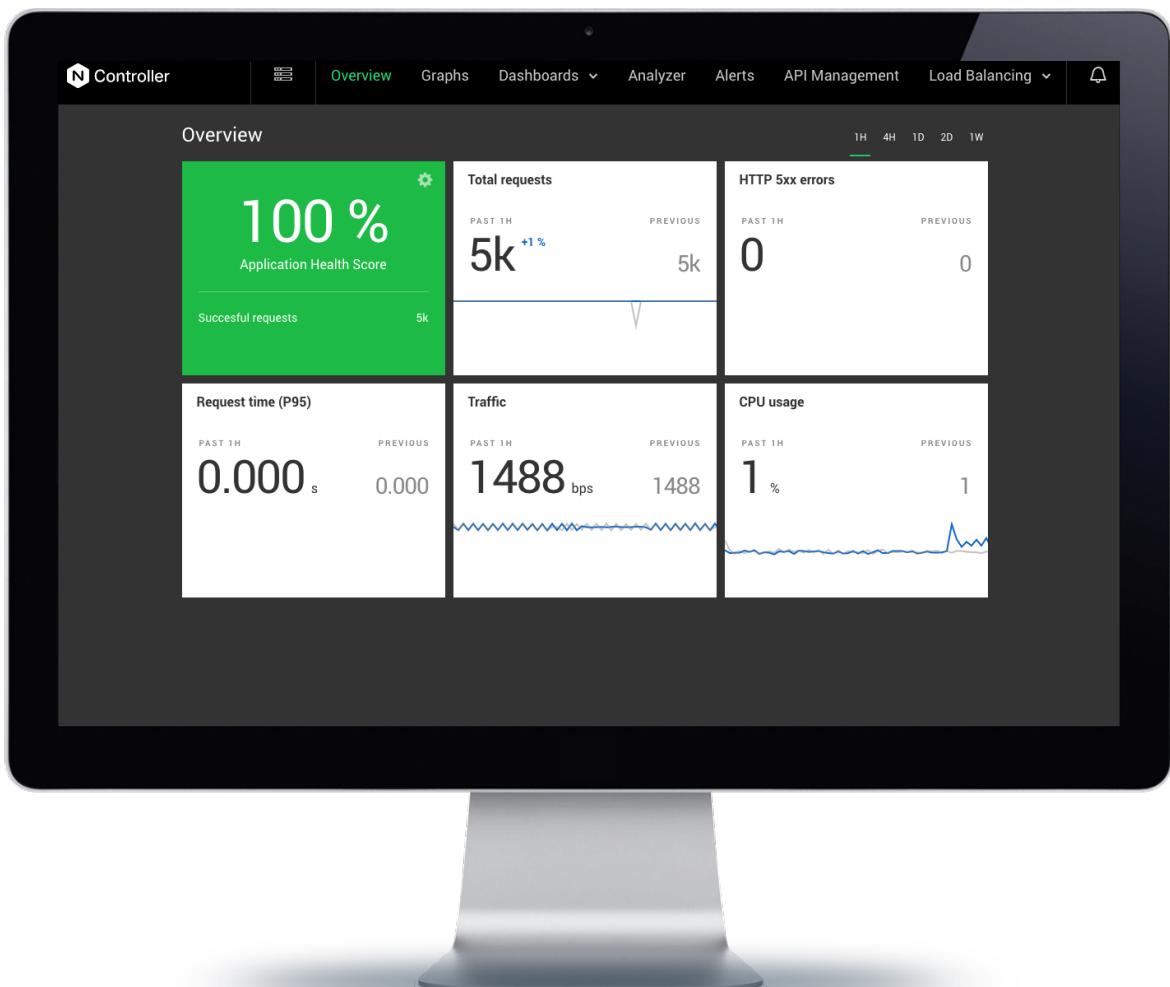
- Load balancer per application
- Load balancer per customer for SaaS providers
- Configuration stored along with application in GitHub
- Fully portable



What is the NGINX Controller?

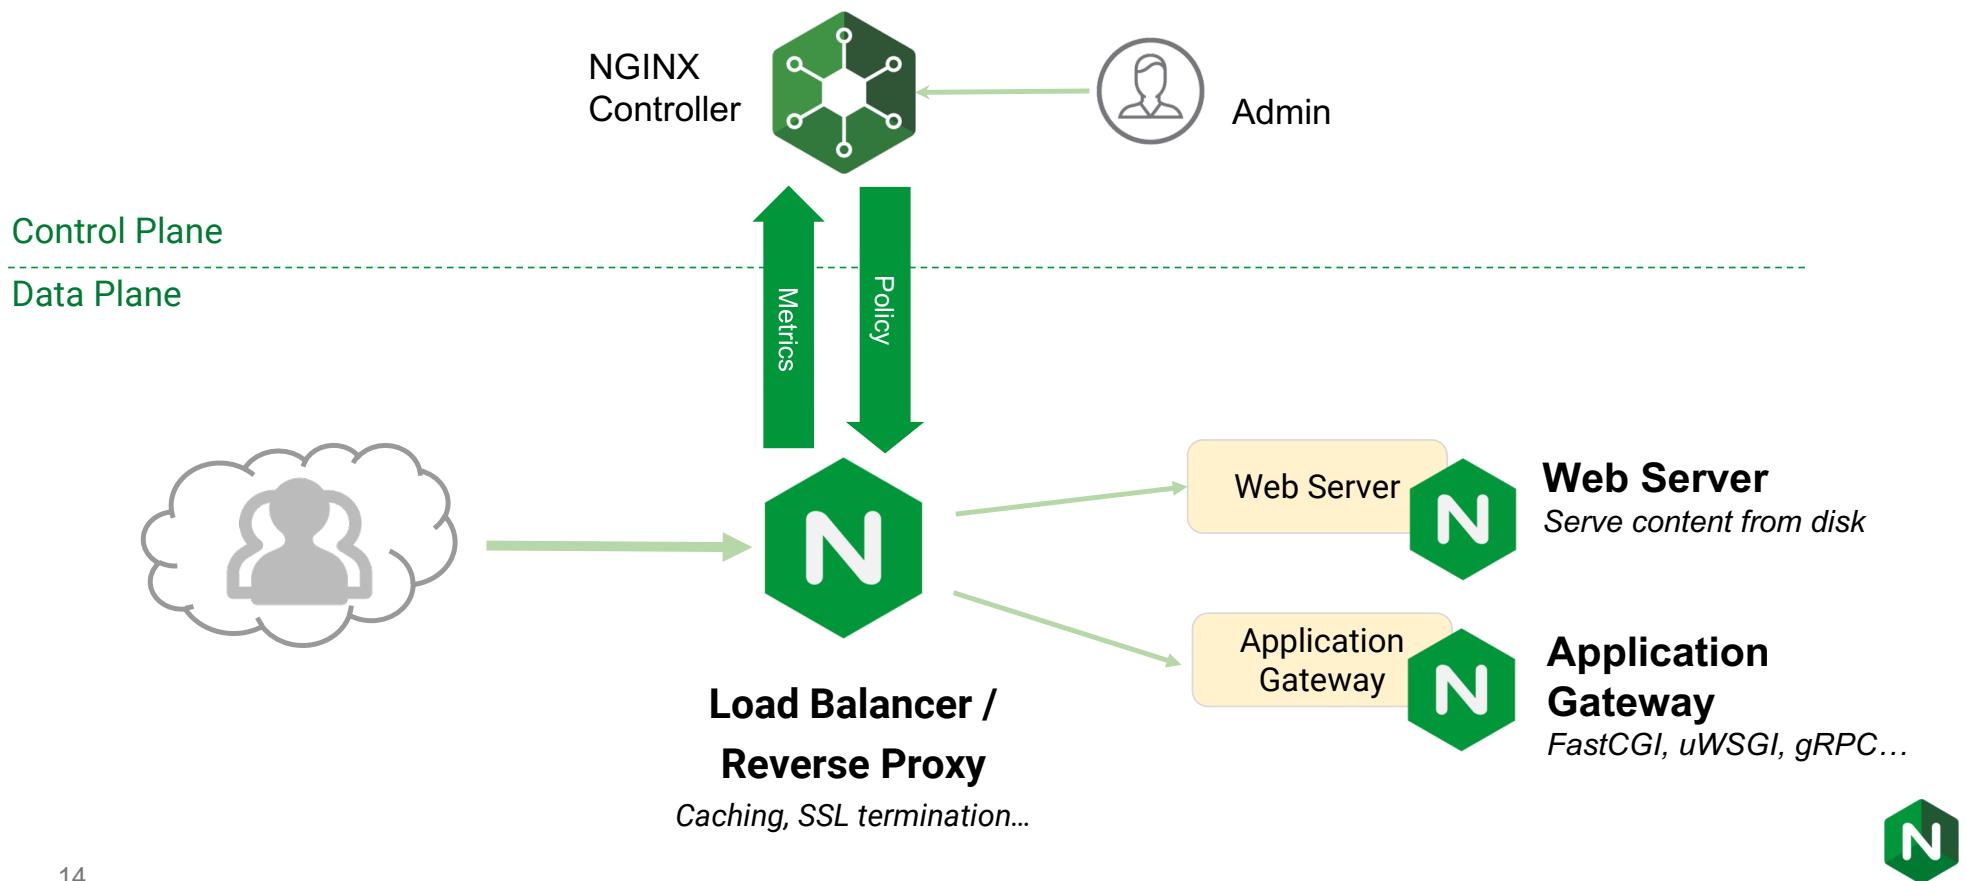
Centralized Monitoring and management

- Alerting
- API management
- Load balancer management
- Configuration analysis
- Customizable dashboards
- Monitoring



What is NGINX and NGINX Controller?

Nginx and the Nginx Controller





Installing NGINX

Nginx Installation Options

- **Official NGINX repo**
 - **Mainline (recommended)** - Actively developed; new minor releases made every 4-6 weeks with new features and enhancements.
 - **Stable** - Updated only when critical issues or security vulnerabilities need to be fixed.
 - **NGINX PLUS** - receives all new features, once they have been tested and proven in NGINX mainline. Additional enterprise-specific features are included in NGINX Plus.
- **OS vendor and other 3rd party repos**
 - Not as frequently updated; e.g. Debian Jessie (8.9) has NGINX 1.6.2
 - Typically built off NGINX mainline branch, sometimes with 3rd party mods
- **Compile from source**
 - Most difficult.-Download the latest version of the NGINX source code, configure, build and install it. You will have the option of building various Nginx module

NGINX Installation: Debian/Ubuntu

Create `/etc/apt/sources.list.d/nginx.list` with the following contents:

```
deb http://nginx.org/packages/mainline/OS/ CODENAME nginx
deb-src http://nginx.org/packages/mainline/OS/ CODENAME nginx
```

- OS – ubuntu or debian depending on your distro
- CODENAME:
 - jessie or stretch for debian
 - trusty, xenial, artful, or bionic for ubuntu

```
$ wget http://nginx.org/keys/nginx_signing.key
$ apt-key add nginx_signing.key
$ apt-get update
$ apt-get install -y nginx
$ /etc/init.d/nginx start
```



NGINX Installation: CentOS/Red Hat

Create `/etc/yum.repos.d/nginx.repo` with the following contents:

```
[nginx]
name=nginx repo
baseurl=http://nginx.org/packages/mainline/OS/OSRELEASE/$basearch/
gpgcheck=0
enabled=1
```

- OS -- rhel or centos depending on your distro
- OSRELEASE -- 6 or 7 for 6.x or 7.x versions, respectively

```
$ yum -y install nginx
$ systemctl enable nginx
$ systemctl start nginx
$ firewall-cmd --permanent --zone=public --add-port=80/tcp
$ firewall-cmd --reload
```



NGINX Plus Installation

Instructions

NGINX Plus packages are available for the following distributions and versions:

- RHEL/CentOS/Oracle Linux
 - 6.5+
 - 7.0+
- Debian
 - 8 (Jessie)
 - 9 (Stretch)
- SLES
 - 12+
- Ubuntu
 - 14.04 (Trusty)
 - 16.04 (Xenial)
 - 17.10 (Artful)
 - 18.04 (Bionic)
- FreeBSD
 - 10.3+
 - 11.0+
- Amazon Linux
 - Amazon Linux 2

To show setup instructions please choose your OS and distribution:

The screenshot shows a dropdown menu titled "Select a distribution". The menu lists several options: RHEL 6/CentOS 6/Oracle Linux 6, RHEL 7.0-7.3/CentOS 7.0-7.3/Oracle Linux 7.0-7.3, RHEL 7.4+/CentOS 7.4+/Oracle Linux 7.4+, Debian, Ubuntu, SLES 12, FreeBSD, Amazon Linux, and Amazon Linux 2. The "RHEL 6/CentOS 6/Oracle Linux 6" option is highlighted with a blue background and white text. Below the dropdown, there is a message: "Information about the latest NGINX Plus updates." At the bottom of the page, there is a navigation bar with links for "Products", "Solutions", "Resources", and "Partners".

- Visit cs.nginx.com/repo_setup
- Select OS from drop down list
- Instructions similar to OSS installation
- Mostly just using a different repo and installing client certificate



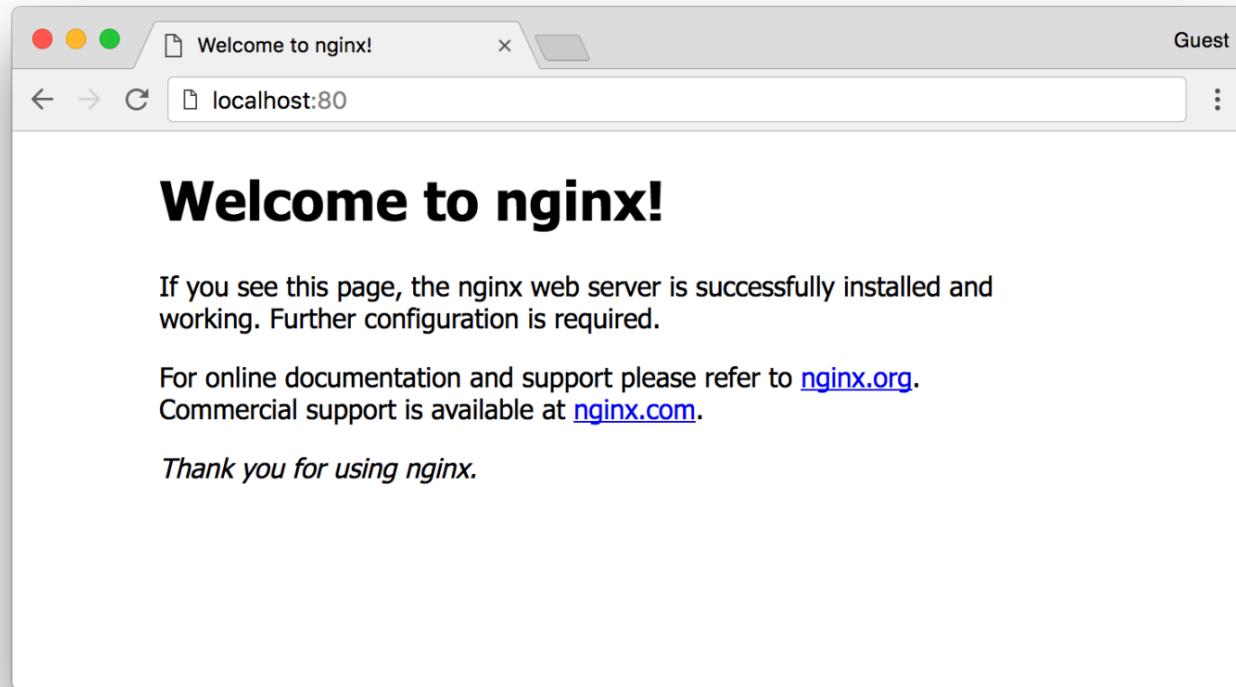
Verifying Installation

```
$ nginx -v
nginx version: nginx version: nginx/1.15.7 (nginx-plus-r17)

$ ps -ef | grep nginx
root      1088      1  0 19:59 ?        00:00:00 nginx: master process /usr/sbin/nginx -c /etc/nginx/nginx.conf
nginx     1092  1088  0 19:59 ?        00:00:00 nginx: worker process
```



Verifying Installation



MORE INFORMATION AT
[NGINX.COM](#)

NGINX Installation Misc

- For more installation details: http://nginx.org/en/linux_packages.html
 - List of all supported distros and CPUs
- For **NGINX Plus**, see: https://cs.nginx.com/repo_setup
 - List of all supported distros and CPUs, including FreeBSD





Essential files, commands and directories

Key NGINX Commands

<code>nginx -h</code>	Shows all command line options
<code>nginx -t</code>	Configuration syntax check
<code>nginx -T</code>	Displays full, concatenated configuration
<code>nginx -V</code>	Shows version and build details
<code>nginx -s reload</code>	Gracefully reload NGINX processes

```
$ sudo nginx -t && sudo nginx -s reload
nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
nginx: configuration file /etc/nginx/nginx.conf test is successful
```

```
$ sudo nginx -T > nginx_support_mm-dd-yy.txt
```

MORE INFORMATION AT
[NGINX.COM](https://nginx.com)

Key System Commands

<code>ps aux grep nginx</code>	To check running processes
<code>ps -ef --forest grep nginx</code>	To check running processes (Show Process Hierarchy in Forest Format)
<code>service nginx status</code> <code>systemctl status nginx</code>	Show Nginx Status
<code>netstat -tulpn</code>	Information and statistics about protocols in use and current TCP/IP network connections.
<code>sudo lsof -i -P -n</code>	Check the listening ports and applications on linux

```
# Path to executable path  
$ /usr/sbin/nginx
```

```
# Default Log Path  
$ /var/log/nginx
```

MORE INFORMATION AT
NGINX.COM

Key Files and Directories

- **/etc/nginx/** # Where all NGINX configuration is stored
- **/etc/nginx/nginx.conf** # Top-level NGINX configuration, should not require much modification
- **/etc/nginx/conf.d/*.conf** # Where your HTTP/S configuration for virtual servers and upstreams goes, e.g. `www.example.com.conf`
- **/etc/nginx/stream.d/*.conf** # Where your TCP/UDP Streams for virtual servers and upstreams goes, e.g. `DNS_53.conf`
- **/var/log/nginx/access.log** # Details about requests and responses
- **/var/log/nginx/error.log** # Details about NGINX errors

Key Files and Directories

/etc/nginx/

nginx.conf

```
#global settings here  
  
http {  
    # HTTP global settings  
    here  
  
    include conf.d/*.conf;  
}
```

Global settings
(tunings, logs, etc)

HTTP block

/etc/nginx/conf.d/

example.com.conf

```
server {  
    listen <parameters>;  
  
    location <url> {  
        -----  
    }  
}  
  
upstream {  
    -----  
}
```

Listen for
requests

Rules to handle
each request

Optional:
upstreams
configurations in
same file,

something.com.conf.disabled Not loaded





Basic configurations

Simple Virtual Server

```
server {  
    listen      80 default_server;  
    server_name www.example.com;  
  
    # ...  
}
```

- `server` defines the context for a virtual server
- `listen` specifies IP/port NGINX should listen on. No IP means bind to all IPs on system
- `server_name` specifies hostname of virtual server



Basic Web Server Configuration

```
server {  
    listen      80 default_server;  
    server_name www.example.com;  
  
    location / {  
        root   /usr/share/nginx/html;  
        index index.html index.htm;  
  
    }  
}
```

- `root` specifies directory where files are stored
- `index` defines files that will be used as an index

- `www.example.com` maps to `/usr/share/nginx/html/index.html` (then `index.htm`)
- `www.example.com/i/file.txt` -> `/usr/share/nginx/html/i/file.txt`



Multiplexing Multiple Sites on One IP

```
server {  
    listen      80 default_server;  
    server_name www.example.com;  
    # ...  
}  
  
server {  
    listen      80;  
    server_name www.example2.com;  
    # ...  
}  
  
server {  
    listen      80;  
    server_name www.example3.com;  
    # ...  
}
```

- NGINX can multiplex a single IP/port using the Host: header.
- default_server defines the virtual server to use if Host header is empty. It is best practice to have a default_server.



Basic SSL Configuration

```
server {  
    listen      80 default_server;  
    server_name www.example.com;  
    return 301 https://$server_name$request_uri;  
}  
  
server {  
    listen 443 ssl default_server;  
    server_name www.example.com;  
    ssl_certificate cert.crt;  
    ssl_certificate_key cert.key;  
    ssl_ciphers HIGH;  
  
    location / {  
        root   /usr/share/nginx/html;  
        index index.html index.htm;  
    }  
}
```

- Force all traffic to SSL is good for security, customer trust and SEO
- Use Let's Encrypt to get free SSL certificates
- Use Mozilla SSL Configuration Generator to generate recommended nginx SSL configurations:
<https://mozilla.github.io/server-side-tls/ssl-config-generator/>



Basic HTTP/2 Configuration

```
server {  
    listen 443 ssl http2 default_server;  
    server_name www.example.com;  
  
    ssl_certificate cert.crt;  
    ssl_certificate_key cert.key;  
}
```

- HTTP/2 improves performance with little to no backend changes
- Add `http2` parameter to `listen` directive of existing SSL-enabled virtual server. HTTP/2 is only supported with SSL in all browsers.
- NGINX only does HTTP/2 client side, server side is still HTTP/1.1. gRPC is a special case.
- Note: HTTP/2 requires OpenSSL 1.0.2 or later to work properly



Basic Reverse Proxy Configuration

```
server {  
    location ~ ^(.+\.php)(.*)$ {  
        fastcgi_split_path_info ^(.+\.php)(.*$);  
  
        # fastcgi_pass 127.0.0.1:9000;  
        fastcgi_pass unix:/var/run/php7.0-fpm.sock;  
  
        fastcgi_index index.php;  
        include fastcgi_params;  
    }  
}
```

- Requires PHP FPM:
`apt-get install -y php7.0-fpm`
- Can also use PHP 5
- Similar directives available for uWSGI and SCGI.
- Additional PHP FPM configuration may be required



Basic Load Balancing Configuration

```
upstream my_upstream {  
    server server1.example.com:80;  
    server server2.example.com:80;  
    least_conn;  
}  
  
server {  
    location / {  
        proxy_set_header Host $host;  
        proxy_pass http://my_upstream;  
    }  
}
```

- `upstream` defines the load balancing pool
- Default load balancing algorithm is round robin. Others available:
 - `least_conn` selects server with least amount of active connections
 - `least_time` factors in connection count and server response time. Available in NGINX Plus only.
- `proxy_pass` links virtual server to upstream
- By default NGINX rewrites Host header to name and port of proxied server. `proxy_set_header` overrides and passes through original client Host header.



Layer 7 Request Routing

```
server {  
    # ...  
  
    location /service1 {  
        proxy_pass http://upstream1;  
    }  
  
    location /service2 {  
        proxy_pass http://upstream2;  
    }  
  
    location /service3 {  
        proxy_pass http://upstream3;  
    }  
}
```

- `location` blocks are used to do Layer 7 routing based on URL
- Regex matching can also be used in `location` blocks



Basic Caching Configuration

```
proxy_cache_path /path/to/cache levels=1:2  
    keys_zone=my_cache:10m max_size=10g  
    inactive=60m use_temp_path=off;  
  
server {  
    location / {  
        proxy_cache my_cache;  
        # proxy_cache_valid 5m;  
        proxy_set_header Host $host;  
        proxy_pass http://my_upstream;  
    }  
}
```

- `proxy_cache_path` defines the parameters of the cache.
- `keys_zone` defines the size of memory to store cache keys in. A 1 MB zone can store data for about 8,000 keys.
- `max_size` sets upper limit of cache size. Optional.
- `inactive` defines how long an object can stay in cache without being accessed. Default is 10 m.
- `proxy_cache` enables caching for the context it is in 



Advanced configurations

Modifications to main nginx.conf

```
user    nginx;  
worker_processes auto;  
  
# ...  
  
http {  
    # ...  
  
    keepalive_timeout 300s;  
    keepalive_requests 100000;  
}
```

- Set in main nginx.conf file
- Default value for worker_processes varies on system and installation source
- auto means to create one worker process per core. This is recommended for most deployments.
- keepalive_timeout controls how long to keep idle connections to clients open. Default: 75s
- keepalive_requests Max requests on a single client connection before its closed.Default: 100
- keepalive_* can also be set per virtual server



HTTP/1.1 Keepalive to Upstreams

```
upstream my_upstream {  
    server server1.example.com;  
    keepalive 32;  
}  
  
server {  
    location / {  
        proxy_set_header Host $host;  
        proxy_http_version 1.1;  
        proxy_set_header Connection "";  
  
        proxy_pass http://my_upstream;  
    }  
}
```

- `keepalive` enables TCP connection cache
- By default NGINX uses HTTP/1.0 with `Connection: Close`
- `proxy_http_version` upgrades connection to HTTP/1.1
- `proxy_set_header` enables keepalive by clearing `Connection: Close` HTTP header



SSL Session Caching

```
server {  
    listen 443 ssl default_server;  
    server_name www.example.com;  
  
    ssl_certificate cert.crt;  
    ssl_certificate_key cert.key;  
  
    ssl_session_cache shared:SSL:10m;  
    ssl_session_timeout 10m;  
}
```

- Improves SSL/TLS performance
- 1 MB session cache can store about 4,000 sessions
- Cache shared across all NGINX workers



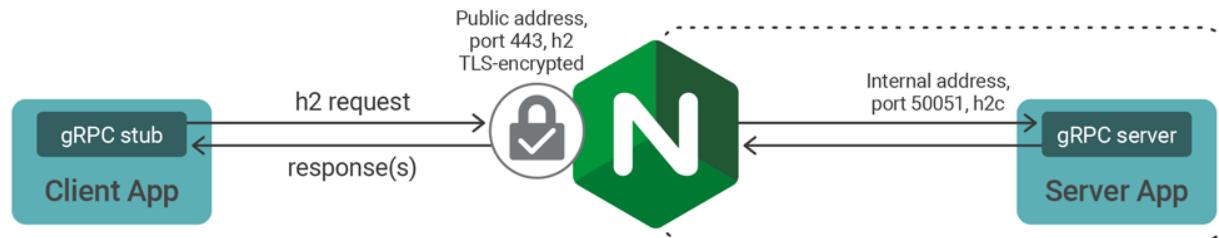
Advanced Caching Configuration

```
proxy_cache_path /path/to/cache levels=1:2  
    keys_zone=my_cache:10m max_size=10g  
    inactive=60m use_temp_path=off;  
  
server {  
    location / {  
        proxy_cache my_cache;  
        proxy_cache_lock on;  
        proxy_cache_revalidate on;  
        proxy_cache_use_stale error timeout updating  
            http_500 http_502 http_503 http_504;  
        proxy_cache_background_update on;  
  
        proxy_set_header Host $host;  
        proxy_pass http://my_upstream;  
    }  
}
```

- `proxy_cache_lock` instructs NGINX to only send one request to the upstream when there are multiple cache misses for the same file.
- `proxy_cache_revalidate` instructs NGINX to use `If-Modified-Since` when refreshing cache.
- `proxy_cache_use_stale` instructs NGINX to serve stale content instead of an error.
- `proxy_cache_background_update` instructs NGINX to do all cache updates in the background. Combined with `proxy_cache_use_stale` updating, stale content will be served.



gRPC Proxying with SSL Termination



```
server {  
    listen 443 ssl http2;  
    ssl_certificate      server.crt;  
    ssl_certificate_key server.key;  
  
    location / {  
        grpc_pass grpc://localhost:50051;  
    }  
}
```

- Configure SSL and HTTP/2 as usual
- Go sample application needs to modified to point to NGINX IP Address and port.



Active Health Checks

```
upstream my_upstream {
    zone my_upstream 64k;
    server server1.example.com slow_start=30s;
    server server2.example.com slow_start=30s;

}

server {
    # ...
    location @health {
        internal;
        health_check interval=5s uri=/test
            match=statusok;
        proxy_set_header HOST www.example.com;
        proxy_pass http://my_upstream;
    }

    match statusok {
        # Used for /test.php health check
        status 200;
        header Content-Type = text/html;
        body ~ "i'm is alive";
    }
}
```

- Polls /test every 5 seconds
- If response is not 200, server marked as failed
- If response body does not contain “I’m alive”, server marked as failed
- Recovered/new servers will slowly ramp up traffic over 30 seconds
- Exclusive to NGINX Plus

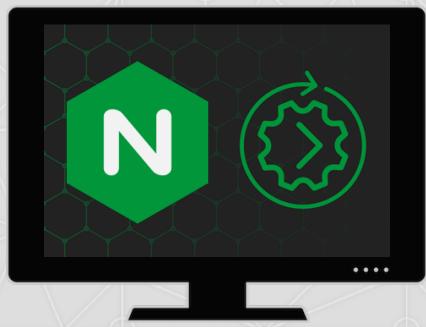


Sticky Cookie Session Persistence

```
upstream my_upstream {  
    server server1.example.com;  
    server server2.example.com;  
  
    sticky cookie name expires=1h  
        domain=.example.com path=/;  
}
```

- NGINX will insert a cookie using the specified *name*
- *expires* defines how long the cookie is valid for. The default is for the cookie to expire at the end of the browser session.
- *domain* specifies the domain the cookie is valid for. If not specified, domain field of cookie is left blank
- *path* specifies the path the cookie is set for. If not specified, path field of cookie is left blank
- Exclusive to NGINX Plus





Monitoring and Logging

NGINX Access Logs

```
access_log /var/log/nginx/access.log;
```

```
192.168.179.1 -- [15/May/2017:16:36:25 -0700] "GET / HTTP/1.1" 200 612 "-" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110 Safari/537.36" "-"
192.168.179.1 -- [15/May/2017:16:36:26 -0700] "GET /favicon.ico HTTP/1.1" 404 571 "http://fmemon-redhat.local/"
"Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110
Safari/537.36" "-"
192.168.179.1 -- [15/May/2017:16:36:31 -0700] "GET /basic_status HTTP/1.1" 200 100 "-" "Mozilla/5.0 (Macintosh; Intel
Mac OS X 10_12_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/58.0.3029.110 Safari/537.36" "-"
```

- Enabled by default. Can be disabled with the `access_log off` directive.
- Nginx uses the **combined log format** (also used by Apache) and includes IP address, date, request, referrer, user agent, etc. You can add additional NGINX variables, e.g. timing and a Log format configurable with the `log_format` directive
- Can enable access logs at a virtual server scope

MORE INFORMATION AT
[NGINX.COM](#)

NGINX Error Logs

```
error_log /var/log/nginx/error.log [level];
```

```
2018/03/22 11:29:08 [error] 12696#12696: upstream timed out (110: Connection timed out) while connecting to upstream,  
health check "" of peer 10.70.88.24:8832 in upstream "Dev.InternalApi"  
2018/03/22 11:29:23 [error] 12696#12696: upstream timed out (110: Connection timed out) while connecting to upstream,  
health check "" of peer 10.70.88.15:8832 in upstream "Dev.InternalApi"  
2018/03/23 15:25:35 [error] 19997#0: *1 open() "/var/www/nginx-default/phpmy-admin/scripts/setup.php" failed (2: No such  
file or directory), client: 80.154.42.54, server: localhost, request: "GET /phpmy-admin/scripts/setup.php HTTP/1.1",  
host: "www.example.com"
```

- Enabled by default. Can be disabled with the `error_log off` directive.
- Can enable access logs at a virtual server scope
- Log to file, stderr, syslog or memory
- Option to log for selected clients

MORE INFORMATION AT
[NGINX.COM](#)

Error Log Levels

```
error_log /var/log/nginx/error.log [level];
```

debug	Detailed Trace
info	General Info
notice	Something Normal
warn	Something Strange
error	Unsuccessful
crit	Important Issue(s)
alert	Fix Now!
emerg	Unusable

MORE INFORMATION AT
NGINX.COM

Extra examples

```
log_format simple escape=json
'{"timestamp": "$time_iso8601", "client": "$remote_addr", "uri": "$uri", "status": "$status"}';

server {
    server_name www.example.com;
    access_log /var/log/nginx/example.log simple;
    error_log syslog:server=192.168.1.1 debug;
}

server {
    server_name www.example2.com;
    map $status $condition {
        ~^2[3] 0;
        default 1;
    }
    access_log /var/log/nginx/example2.log simple custom if=$condition;
    error_log /var/log/nginx/example2_error.log info;
}
```

Example log parsing commands:

<code>tail -f 10 error.log</code>	Tail error logs (last 10 lines)
<code>tail -f 10 access.log grep 127.0.0.1</code>	Tail and grep (filter) access logs
<code>cat access.log cut -d '"' -f3 cut -d ' ' -f2 sort uniq -c sort -rn</code>	Sort access by Response Codes
<code>awk '(\$9 ~ /404/)' access.log awk '{print \$7}' sort uniq -c sort -rn</code>	Which links are broken (HTTP 404)?
<code>awk -F\" '{print \$2}' access.log awk '{print \$2}' sort uniq -c sort -r</code>	What are my most requested links?

MORE INFORMATION AT
[NGINX.COM](https://nginx.com)

NGINX Stub Status Module

```
server {  
    location /basic_status {  
        stub_status;  
    }  
}
```

- Provides aggregated NGINX statistics
- Restrict access so it's not publicly visible

```
$ curl http://127.0.0.1/basic_status  
Active connections: 1  
server accepts handled requests  
 7 7 7  
Reading: 0 Writing: 1 Waiting: 0
```

MORE INFORMATION AT
NGINX.COM

NGINX Plus Extended Status Module

The screenshot shows the NGINX Plus Extended Status Module interface. At the top, there's a navigation bar with tabs: Server zones (green checkmark), Upstreams (red X), TCP/UDP Zones (green checkmark), TCP/UDP Upstreams (green checkmark), Caches (yellow exclamation mark), Shared zones (grey), and a gear icon. Below the navigation bar, there are several sections displaying real-time metrics:

- Connections**: Current: 46, Accepted/s: 6, Active: 2, Idle: 44, Dropped: 0, Accepted: 87645398.
- Requests**: Current: 1, Req/s: 47, Total: 241744717.
- Server zones**: Total: 3, Problems: 0.
- Upstreams**: Total: 4, Alerts: 2.
- TCP/UDP Zones**: Conn total: 338228, Conn current: 0, Conn/s: 0.
- TCP/UDP Upstreams**: Total: 3, Problems: 0.
- Caches**: Total: 1, Warnings: 1.

On the left side, there are traffic details: Traffic In: 346 B/s, Out: 5.11 kB/s; and a servers section: All: 7 / Up: 5, Failed: 2.

```
upstream my_upstream {  
    #...  
}  
  
zone my_upstream 64k;  
  
  
server {  
    #...  
    status_zone my_virtual_server;  
}
```

- Provides detailed NGINX Plus statistics
- Over 100+ additional metrics
- Monitoring GUI also available; see demo.nginx.com
- Exclusive to NGINX Plus

MORE INFORMATION AT
NGINX.COM

NGINX Plus Dashboard

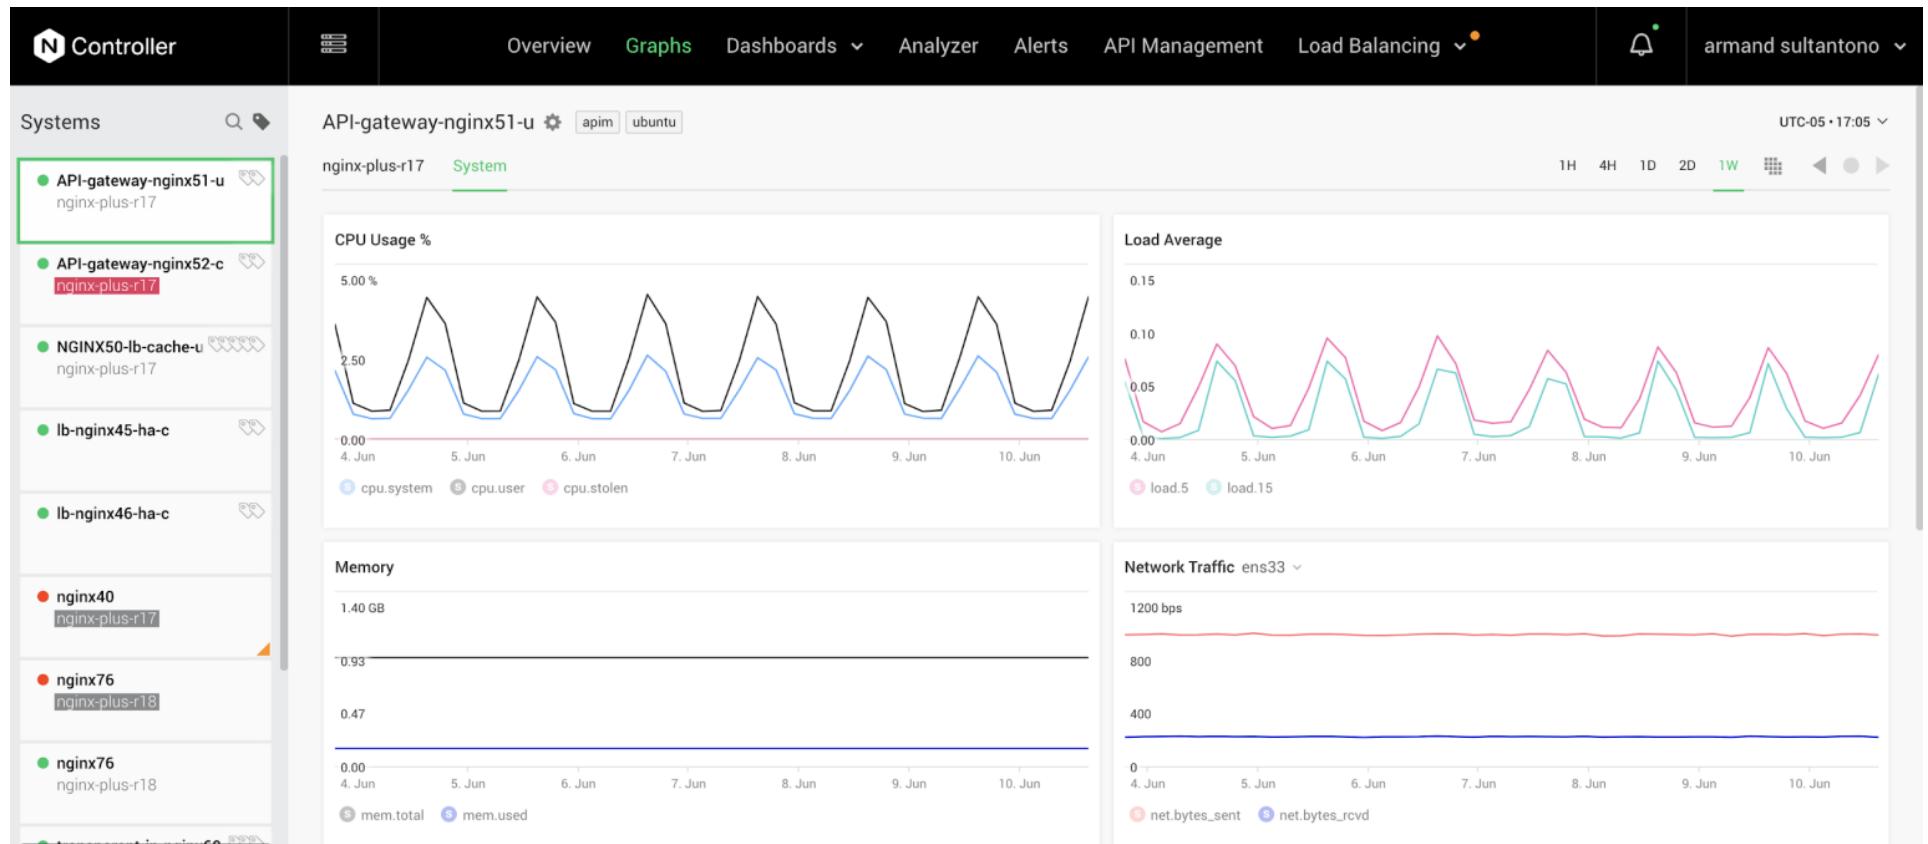
The dashboard displays the following key information:

- nginx-plus-r12-p2 (1.11.10)**
- Address**: 206.251.255.64
- PID**: 98240
- Uptime**: 2d 7h 56m
- Connections**: Accepted: 1911755
- Current**: 32
- Accepted/s**: 3
- Active**: 3
- Idle**: 29
- Dropped**: 0
- Server zones**: Total 3 / Problems 0
- Upstreams**: Total 4 / Alerts 2
- TCP/UDP Zones**: Conn total: 315662, Conn current: 0, Conn/s: 1
- TCP/UDP Upstreams**: Total 3 / Problems 0
- Traffic**: In: 93.0 B/s, Out: 4.33 kB/s
- Servers**: All: 7 / Up: 4, Failed: 3
- Traffic**: In: 107 B/s, Out: 14.1 kB/s
- Servers**: All: 11 / Up: 5, Failed: 0

```
"nginx_build": "nginx-plus-r12-p2",
"nginx_version": "1.11.10",
"pid": 98240,
"ppid": 50622,
"processes": {
    "respawned": 0
},
"requests": {
    "current": 1,
    "total": 9915307
},
"server_zones": {
    "hg.nginx.org": {
        "discarded": 9150,
        "processing": 0,
        "received": 146131844,
        "requests": 597471,
        "responses": {
            "1xx": 0,
            "2xx": 561986,
            "3xx": 12839,
            "4xx": 7081,
            "5xx": 6415,
            "total": 588321
        },
        "sent": 14036626711
},
```

- Over [0 metrics](#) compared to open source NGINX
- Per virtual server and per backend server statistics
- JSON output to export to your favorite monitoring tool
- See demo.nginx.com for live demo

NGINX Controller



NGINX.COM



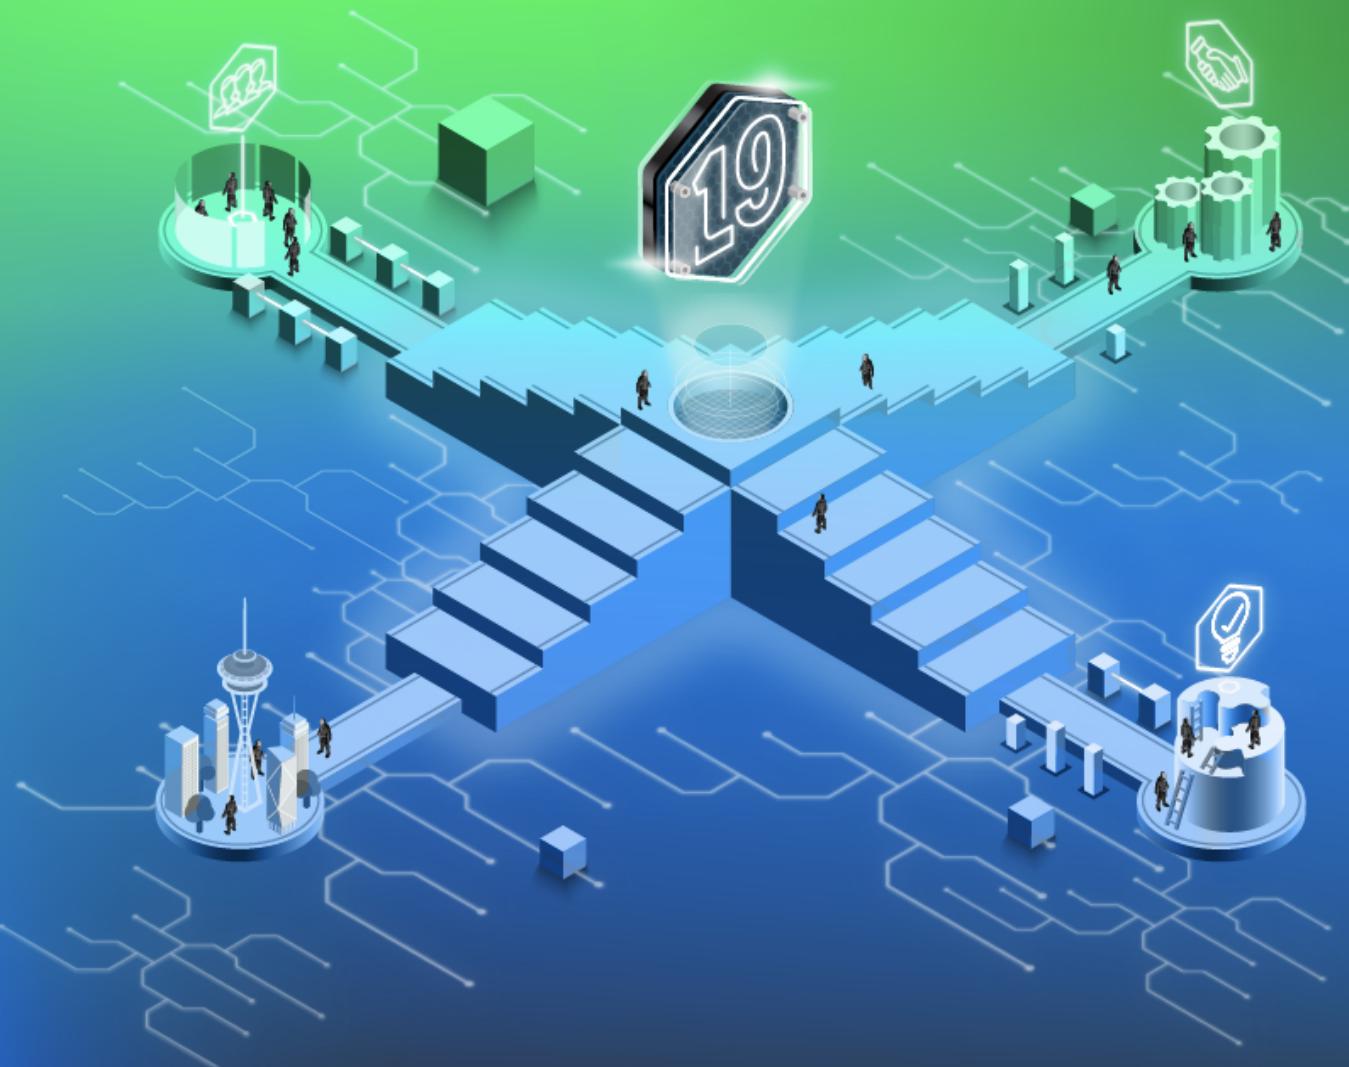
Summary

Summary

- It is recommended to use the NGINX mainline branch for most deployments
- All configuration should go into separate files in /etc/nginx/conf.d/*.conf
- Forcing all traffic to SSL improves security and improves search rankings
- Keepalive connections improve performance by reusing TCP connections
- SSL session caching and HTTP/2 improve SSL performance
- NGINX status module and logging capability provide visibility
- NGINX Plus is recommended for all production, load balancing, API gateway deployments

Try NGINX Plus for free at nginx.com/free-trial-request





NGINX Conf 2019

Sep. 9–12, 2019 | Seattle, WA

The official event
for all things NGINX

REGISTER NOW:
www.nginx.com/nginxconf/2019

Q & A

Try NGINX Plus free for 30 days: nginx.com/free-trial-request

FAQs

- Is NGINX Plus a software package or a “virtual appliance” (an image of a virtual machine for ESX/Xen/etc.)?
- What kind of 3rd party software is bundled along with NGINX Plus?
- Does NGINX Plus gather user details or installation details, or phone home?
- Can we leave everything as-is after we successfully do a trial of NGINX Plus and have purchased subscriptions, or will we need to re-install/re-deploy NGINX Plus?
- What happens when my NGINX Plus subscription expires?

MORE INFORMATION AT
NGINX.COM